

Review

- Last Class:
 - Review of instruction set architecture (ISA)
 - Instruction format
 - Addressing mode
- · Today's class:
 - Review of pipeline
 - Single cycle processor → CCT problem
 - Pipelined execution → improved ILP
 - Hazards hurt performance
 - Multi-cycle pipelines
- Announcement and reminder
 - Office hour starts today (virtual, zoom link on syllabus)

(1)

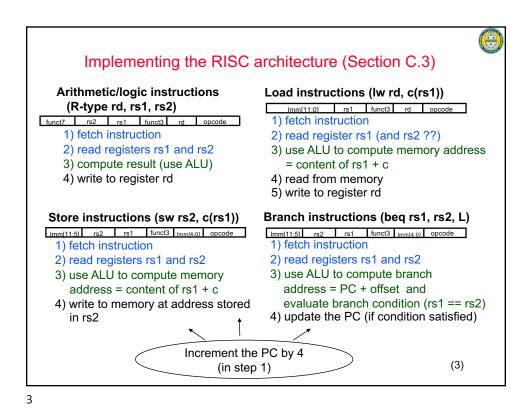
1

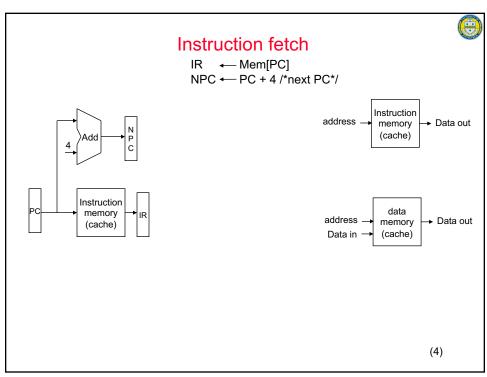
Computer Pipelines

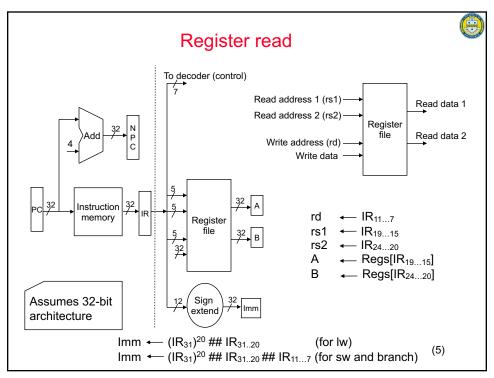


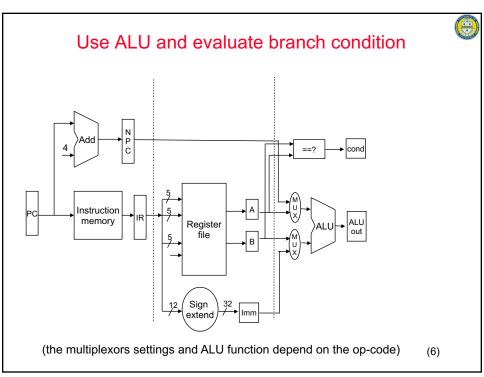
- · RISC features that facilitates efficient pipelining:
 - all instructions have the same length,
 - registers located in the same place in instruction format,
 - memory operands only in loads or stores
- We will first review a non-pipelined RISC architecture that executes:
- Review you should know about:
 - multiplexors,
 - register files,
 - ALU's
 - arithmetic Vs logical shifts
 - program counter (PC), status word (PS) and instruction register (IR),
 - Memory address registers (MAR) and memory data registers (MDR)
 - combinational Vs sequential circuits

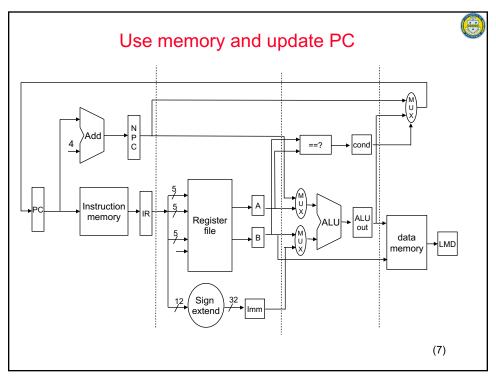
(2)

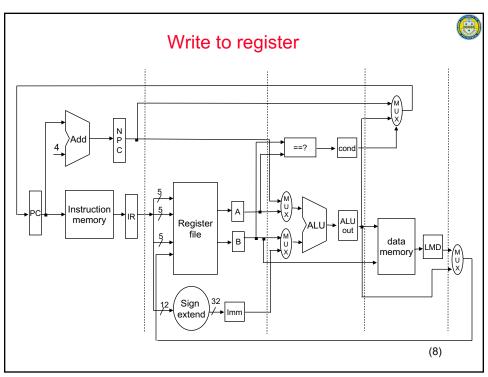


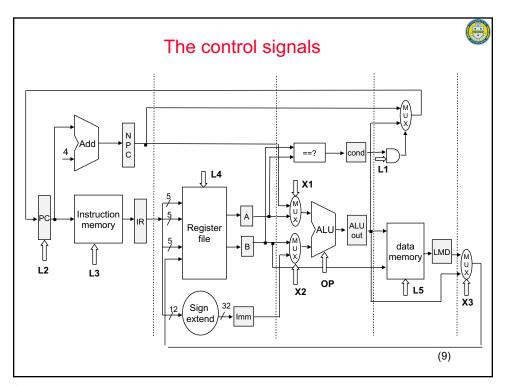






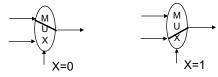






The control signals

• X1, X2, X3 = select multiplexor input (one bit each)



- L1 = set if the instruction is a branch (one bit)
- L2 = loads the PC (one bit)
- L3 = read the instruction memory (one bit)
- L4 = write register (one bits) 0 = no write, 1 = write
- L5 = read/write the data memory (two bits)
 - 00 = no-op, 01 = read, 10 = write
- OP = ALU control (?? Bits) 0..0 = no-op, 1..1 = add, 10.. = others

(10)



Control Signal Table

Operation	RegDst	RegWrite	ALUSrc	ALUOp	MemWrite	MemRead	MemToReg
add	1	1	0	010	0	0	0
sub	1	1	0	110	0	0	0
and	1	1	0	000	0	0	0
or	1	1	0	001	0	0	0
slt	1	1	0	111	0	0	0
lw	0	1	1	010	0	1	1
sw	Х	0	1	010	1	0	Χ
beq	X	0	0	110	0	0	X

- sw and beq are the only instructions that do not write any registers.
- Iw and sw are the only instructions that use the constant field. They
 also depend on the ALU to compute the effective memory address.
- · ALUOp for R-type instructions depends on the instructions' func field.
- The PCSrc control signal (not listed) should be set if the instruction is beq *and* the ALU's Zero output is true.

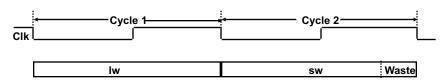
(11)

11

Single Cycle Disadvantages & Advantages



- Uses the clock cycle inefficiently the clock cycle must be timed to accommodate the slowest instruction
 - especially problematic for more complex instructions like floating point (FP) multiply



 May be wasteful of area since some functional units (e.g., adders) must be duplicated since they can not be shared during a clock cycle

but

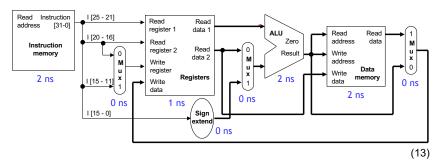
Single cycle datapath is simple and easy to understand

(12)

CCT Problem

- If all instructions must complete within one clock cycle, then the cycle time has to be large enough to accommodate the slowest instruction.
- For example, lw \$t0, -4(\$sp) needs 8ns, assuming the delays shown here.

reading the instruction memory 2ns reading the base register \$sp 1ns computing memory address \$sp-4 2ns 8ns reading the data memory 2ns storing data back to \$t0 1ns

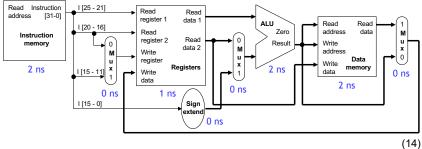


13

CCT Problem

- If we make the cycle time 8ns then every instruction will take 8ns, even if they don't need that much time.
- For example, the instruction add \$s4, \$t1, \$t2 really needs just 6ns.

reading the instruction memory 2ns reading registers \$t1 and \$t2 1ns 6ns computing \$t1 + \$t2 2ns storing the result into \$s0 1ns



Another Example



- With these same component delays, a sw instruction would need 7ns, and beq would need just 5ns.
- Let's consider the gcc instruction mix:

Instruction	Frequency
Arithmetic	48%
Loads	22%
Stores	11%
Branches	19%

- With a single-cycle datapath, each instruction would require 8ns.
- But if we could execute instructions as fast as possible, the average time per instruction for gcc would be:

$$(48\% \times 6ns) + (22\% \times 8ns) + (11\% \times 7ns) + (19\% \times 5ns) = 6.36ns$$

· The single-cycle datapath is about 1.26 times slower!

(15)

15

Pipelining: Its Natural!



· Car wash Example



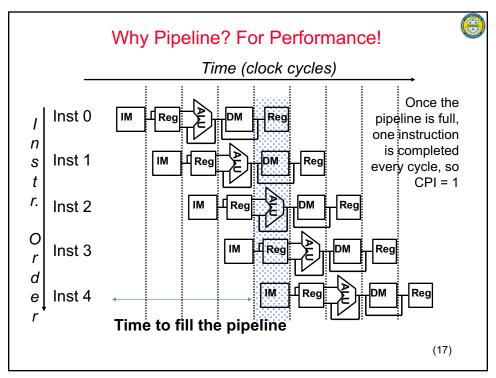


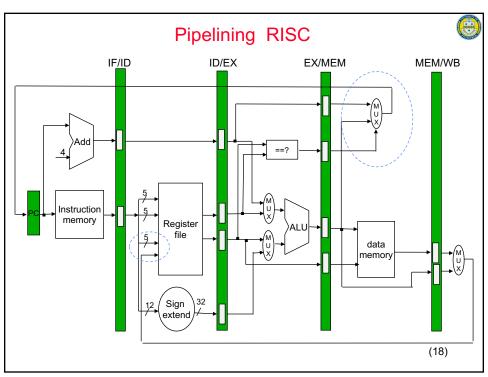


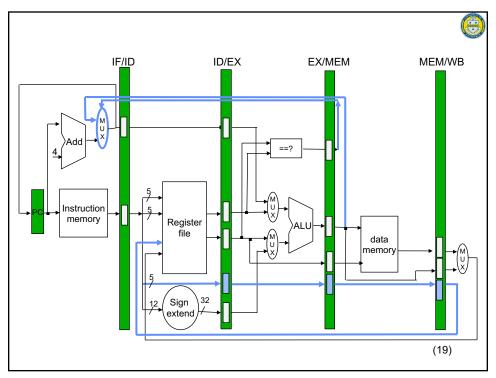


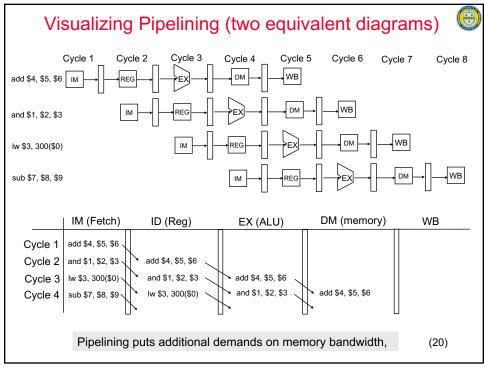
- Pipelining doesn't help latency of single task, it helps throughput of entire workload
- Multiple tasks operating simultaneously
- Potential speedup = Number pipe stages
- Pipeline rate limited by slowest pipeline stage
- Unbalanced lengths of pipe stages reduces speedup
- Time to "fill" pipeline and time to "drain" it reduces speedup

(16)











Limits to pipelining (Sections C.1 and C.2)

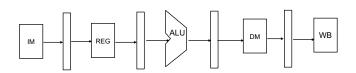
- Hazards prevent next instruction from executing during its designated clock cycle
 - Structural hazards: HW cannot support some combination of instructions.
 - Data hazards: An instruction depends on result of prior instruction still in the pipeline.
 - Control hazards: Pipelining of branches & other instructions stall the pipeline until the hazard bubbles in the pipeline
- · Dependencies backward in time cause hazards
- · Can usually resolve hazards by waiting
 - pipeline control must detect the hazard
 - and take action to resolve hazards

(21)

21

Structural Hazards





Potential problem: Both REG and WB use the register file

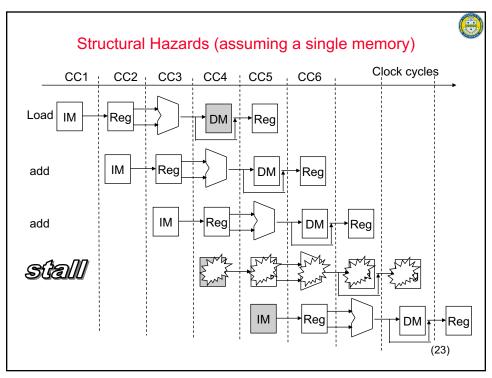
Solution : Read from register file during the first half of a

cycle and write to register file during the second

half of a cycle.

Potential problem: Both IM and DM use memory **Solution**: use separate memories (caches)

(22)



Speed Up Equation for Pipelining

 $\mathit{CPI}_{\mathit{pipelined}}$ = Ideal CPI + stall cycles per instruction

$$Speedup = \frac{CPI_{unpipelined}}{CPI_{pipelined}} \mathbf{x} \frac{Clock Cycle_{unpipelined}}{Clock Cycle_{pipelined}}$$

Example:

- · Machine A: pipelined (with some depth) and dual ported memory
- Machine B: pipelined (same depth as A), but single ported memory, and a 1.05 times faster clock rate
- Ideal CPI = 1 for both, and lw/sw are 40% of instructions executed

(24)

Three types hazards caused by Data Dependence



Read After Write (RAW)
 Instr_{i+k} reads operand before Instr_i writes it

Write After Read (WAR)

Instr_{i+k} writes operand *before Instr_i* reads it

- Gets wrong operand
- Can't happen in RISC 5 stage pipeline (why?)

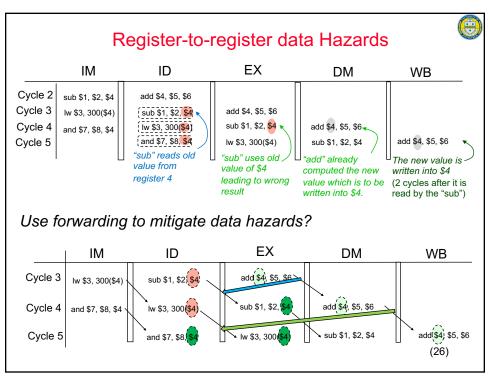


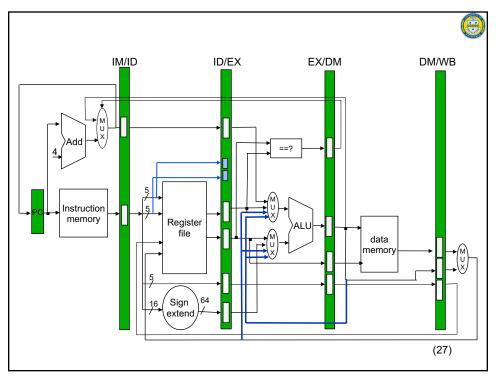
Write After Write (WAW)

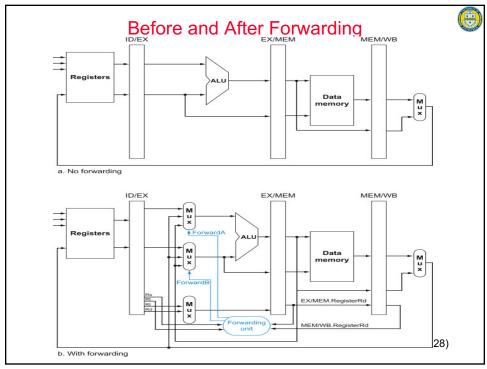
Instr_{i+k} writes operand *before Instr_i* writes it

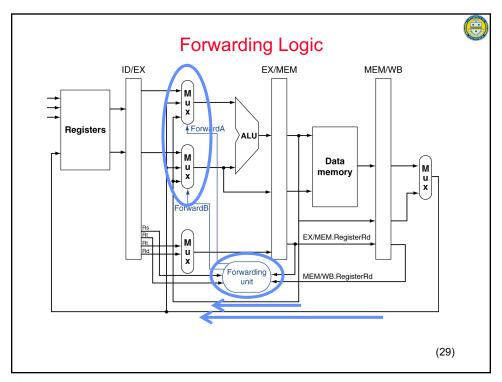
- Leaves wrong result
- Can't happen in RISC 5 stage pipeline (why?)
- Forwarding mitigates data hazards (load-use hazard?)
- · Will see WAR and WAW in later more complicated pipes

(25)









Control Values for the Forwarding Multiplexors



Mux control	Source	Explanation
ForwardA = 00	ID/EX	The first ALU operand comes from the register file.
ForwardA = 10	EX/MEM	The first ALU operand is forwarded from the prior ALU result.
ForwardA = 01	MEM/WB	The first ALU operand is forwarded from data memory or an earlier ALU result.
ForwardB = 00	ID/EX	The second ALU operand comes from the register file.
ForwardB = 10	EX/MEM	The second ALU operand is forwarded from the prior ALU result.
ForwardB = 01	MEM/WB	The second ALU operand is forwarded from data memory or an earlier ALU result.

(30)



Data Forwarding Control Conditions

1. EX/MEM Forward Unit:

```
if (EX/MEM.RegWrite
and (EX/MEM.RegisterRd != 0)
and (EX/MEM.RegisterRd = ID/EX.RegisterRs))
      ForwardA = 10
if (EX/MEM.RegWrite
and (EX/MEM.RegisterRd != 0)
and (EX/MEM.RegisterRd = ID/EX.RegisterRt))
      ForwardB = 10
```

Forwards the result from the previous instr. to either input of the ALU

MEM/WB Forward Unit:

```
if (MEM/WB.RegWrite
and (MEM/WB.RegisterRd != 0)
and (MEM/WB.RegisterRd = ID/EX.RegisterRs))
      ForwardA = 01
if (MEM/WB.RegWrite
and (MEM/WB.RegisterRd != 0)
and (MEM/WB.RegisterRd = ID/EX.RegisterRt))
      ForwardB = 01
```

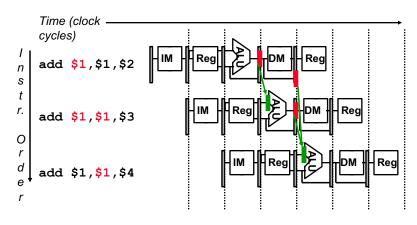
Forwards the result from the second previous instr. to either input of the ALU

(31)

31

Yet Another Complication!

Another potential data hazard can occur when there is a conflict between the result of the WB stage instruction and the MEM stage instruction – which should be forwarded?



(32)

Corrected Data Forwarding Control Conditions 1. EX/MEM Forward Unit:



if (EX/MEM.RegWrite

Forwards the result from the previous instr. to either input of the ALU

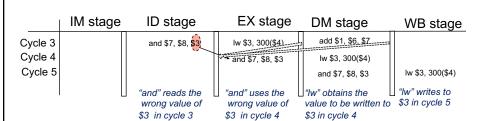
MEM/WB Forward Unit:

```
if (MEM/WB.RegWrite
and (MEM/WB.RegisterRd != 0)
and !(EX/MEM.RegWrite and (EX/MEM RegisterRD != 0)
   and (EX/MEM.RegisterRd = ID/EX.RegisterRs))
and (MEM/WB.RegisterRd = ID/EX.RegisterRs))
      ForwardA = 01
                                       Forwards the result from
                                       the previous or second
                                       previous instr. to either
if (MEM/WB.RegWrite
                                       input of the ALU
and (MEM/WB.RegisterRd != 0)
and !(EX/MEM.RegWrite and (EX/MEM.RegisterRd != 0)
   and (EX/MEM.RegisterRd = ID/EX.RegisterRt)
and (MEM/WB.RegisterRd = ID/EX.RegisterRt))
      ForwardB = 01
                                                     (33)
```

33

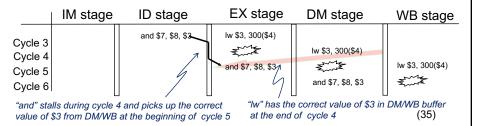
Forwarding may not be enough





Problem: can't use forwarding at the beginning of cycle 4 since "lw" produces the data to be written in \$3 at the end of cycle 4

Solution: need to combine forwarding with stalling the pipe.





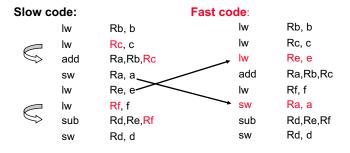
Software Scheduling to Avoid Load Hazards

Try producing fast code for

$$a = b + c$$
;

$$d = e - f$$
;

assuming a, b, c, d ,e, and f in memory.



(36)