



## Review

- Last Class:
  - Tomasulo speculative OOO execution
  - Multi-issue OOO execution
- Today's class:
  - Multi-threading
  - SMT
  - Review of memory hierarchy
- Announcement and reminder
  - HW2 will be distributed on Canvas tonight

(149)

149



## Dynamic scheduling (§3.8)


- Extends Tomasulo's algorithm to issue two (or more) instructions simultaneously to reservation stations.
- Either issue an instruction every half clock cycle, or double the logic to handle two instructions at once.
- Use the same logic. Some restrictions may be used to simplify hardware.
  - Example: issue only one FP and one int. operation every clock cycle. This reduces the load on the register files.
  - Do not issue dependent instructions in the same cycle
- To deal with control hazards without speculation (no ROBs): instructions following a branch can be issued but cannot start execution before the branch is resolved.
- Will look at the execution of

```
L1: fld    F0, 0(R1)
    fadd.d  F4, F0, F2
    fsd     F4, 0(R1)
    addi    R1, R1, -8
    bne     R1, R2, L1
```

(150)

150

Dual issue with one int. and one FP add unit (not pipelined). Latency = 1, 2 and 3 for int. operations, loads and FP adds (one CDB).




Iteration	Instruction	Issued at	Executes	Mem access	Write CDB	comments
1	fld F0, 0(R1)	1	2	3	4	First issue
1	fadd.d F4, F0, F2	1	5	8	8	Wait for fld
1	fsd F4, 0(R1)	2	3	9		Wait for fadd.d
1	addi R1, R1, -8	2	4		5	Wait for ALU
1	bne R1, R2, L1	3	6			wait for addi
2	fld F0, 0(R1)	4	7	8	9	Wait for bne
2	fadd.d F4, F0, F2	4	10	13	13	Wait for fld
2	fsd F4, 0(R1)	5	8	14		Wait for fadd.d
2	addi R1, R1, -8	5	9		10	Wait for ALU
2	bne R1, R2, L1	6	11			wait for addi
3	fld F0, 0(R1)	7	12	13	14	Wait for bne
3	fadd.d F4, F0, F2	7	15	18	18	Wait for fld
3	fsd F4, 0(R1)	8	13	19		Wait for fadd.d
3	addi R1, R1, -8	8	14		15	Wait for ALU
3	bne R1, R2, L1	9	16			wait for addi

- No speculation: Instruction after *bne* cannot be issued in same cycle since the branch is not yet predicted (orange arrows)
  - *fld*, *fsd*, *addi* and *bne* use the same integer Execute unit (see blue circles)
  - Assume as many reservation stations as needed
- (151)

151

If we have a separate int. unit for memory address calculation and we have a second CDB.



Iteration	Instruction	Issued at	Executes	Mem access	Write CDB	comments
1	fld F0, 0(R1)	1	2	3	4	First issue
1	fadd.d F4, F0, F2	1	5	8	8	Wait for fld
1	fsd F4, 0(R1)	2	3	9		Wait for fadd.d
1	addi R1, R1, -8	2	3		4	Exec. earlier
1	bne R1, R2, L1	3	5			wait for addi
2	fld F0, 0(R1)	4	6	7	8	Wait for bne
2	fadd.d F4, F0, F2	4	9	12	12	Wait for fld
2	fsd F4, 0(R1)	5	7	13		Wait for fadd.d
2	addi R1, R1, -8	5	6		7	Exec. earlier
2	bne R1, R2, L1	6	8			wait for addi
3	fld F0, 0(R1)	7	9	10	11	Wait for bne
3	fadd.d F4, F0, F2	7	12	15	15	Wait for fld
3	fsd F4, 0(R1)	8	10	16		Wait for fadd.d
3	addi R1, R1, -8	8	9		10	Exec. earlier
3	bne R1, R2, L1	9	11			wait for addi


No speculation: no instruction after *bne* can start before *bne* completes execution (cycle 5)

(152)

152

## Multiple issue with speculation

Consider same example (as last slide)



Iteration	Instruction	Issued at	Executes	Mem access	Write CDB	comments
1	fld F0, 0(R1)	1	2	3	4	First issue
1	fadd.d F4, F0, F2	1	5	8		Wait for fld
1	fsd F4, 0(R1)	2	3	9		Wait for fadd.d
1	addi R1, R1, -8	2	3		4	
1	bne R1, R2, L1	3	5			wait for addi
2	fld F0, 0(R1)	4	3	4	5	No wait for bne
2	fadd.d F4, F0, F2	4	7	12	10	Wait for fld
2	fsd F4, 0(R1)	5	4	5	13	Wait for fadd.d
2	addi R1, R1, -8	5	6		7	
2	bne R1, R2, L1	6	5	8		wait for addi
3	fld F0, 0(R1)	7	6	7	8	No wait for bne
3	fadd.d F4, F0, F2	7	6	12	10	Wait for fld
3	fsd F4, 0(R1)	8	7	10	8	Wait for fadd.d
3	addi R1, R1, -8	8	7	8	10	
3	bne R1, R2, L1	9	8	11	10	wait for addi

Allow speculation, and may issue a branch with another instruction

(153)

153

## Limits to Multi-Issue Machines

- Latencies of units: many operations must be scheduled
- Need about (Pipeline Depth x No. Functional Units) of independent instructions.
- Need More instruction fetch bandwidth (easy)
- Need Duplicate FUs to get parallel execution (easy)
- Need more ports in Register File (hard) and more memory bandwidth (harder)
- Impact of decoding multiple instructions on clock rate and pipeline depth.
- Decode issue in Superscalar: how wide is practical?
- More logic leads to larger power consumption → lower performance per watt.

(154)

154



## Return address prediction

- Most unconditional branches come from function returns
- The same procedure can be called from multiple sites
  - Causes the buffer to potentially forget about the return address from previous calls

## Integrated instruction Fetch unit that performs

- Branch prediction
- Instruction prefetch

## Energy Efficiency

- Speculation is only energy efficient when it significantly improves performance

## Value Prediction

- Loads that load from a constant pool
- Instruction that produces a value from a small set of values
- Not incorporated into modern processors

(155)

155

## Dynamically scheduling memory ops

- ❑ Compilers must schedule memory ops conservatively
- ❑ Options for hardware:
  1. Don't execute any load until all prior stores execute (**conservative**)
  2. Execute loads as soon as possible, detect violations (**aggressive**)
    - When a store executes, it checks if any later loads executed too early (to the same address). If so, flush pipeline and restart
  3. Learn violations over time, selectively reorder (**predictive**)

### Conservative

```
ld r2,4(sp)
ld r3,8(sp)
add r3,r2→r1
st r1,0(sp)
ld r5,0(r8) //stall
ld r6,4(r8)
sub r5,r6→r4
st r4,8(r8)
```

### Aggressive (Wrong Outcome?)

```
ld r2,4(sp)
ld r3,8(sp)
ld r5,0(r8) //does r8==sp?
add r3,r2→r1
ld r6,4(r8) //does r8+4==sp?
st r1,0(sp)
sub r5,r6→r4
st r4,8(r8)
```

(156)

156

## Store Queue with load bypassing

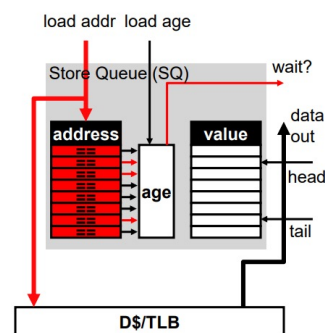
- ❑ Loads can be allowed to bypass stores (if no aliasing)
  - Requires checking addresses of older stores
  - Addresses of older stores must be known in order to check
- ❑ To implement, use separate load queue (LQ) and store queue (SQ)
  - Think of separate RS for loads and stores
- ❑ Need to know the relative order of instructions in the queues – “Age”: new field added to both queues
  - Age represents position of load/store in the program
  - A simple counter incremented during the in-order dispatch

(157)

157

## Store Queue with load bypassing

- ❑ Loads: for the oldest ready load in LQ, check the “Addr” of older stores in SQ
  - If any store with an uncomputed or matching “Addr”, load cannot issue
  - Check SQ in parallel with accessing D\$
- ❑ • Stores: can always execute when at ROB head



(158)

158

## Memory forwarding


- ❑ Stores write cache at Commit (until then hold store value and address in a Store Queue)
  - Commit is in-order, so allows stores to (also) be “undone” on branch mis-predictions, etc.
- ❑ Loads read cache
  - Early execution of loads is critical
- ❑ Forwarding
  - Allow store to (following) load communication before store Commits
  - Conceptually like register forwarding, but different implementation
    - Why? Addresses unknown until Execute

(159)

159

## Store Queue with load forwarding

- ❑ Store Queue (SQ)
  - Holds all in-flight stores
  - CAM: searchable by address
  - Age logic: determine *youngest* matching store that is *older* than the load
- ❑ Store execution
  - Write into SQ
    - address + data value
  - Store from SQ into D\$ at Commit
- ❑ Load execution
  - Assoc search SQ addresses
    - Match? Forward youngest matching store value to load request
  - Else read value from D\$



(160)

160

## Load speculation

- ❑ Allow load to proceed when there are undetermined store addresses in the SQ.
- ❑ Speculation requires two things ...
  1. Detection of mis-speculations (guessed that memory addresses didn't match and it turns out that they did)
    - How can we do this?
  2. Recovery from mis-speculations
    - Squash instr's after offending load (they may be using the load data)
    - Saw how to squash instr's after mispredicated branches: same method

(161)

161

## LSQ = Store Queue + Load Queue

- ❑ **Store Queue:** handles forwarding
  - Written into by stores (@ store execute)
  - Searched by loads (@ load execute) for store forwarding
  - Write to data cache (@ store Commit)
- ❑ **Load Queue:** detects ordering violations
  - Written into by loads (@ load execute)
  - Searched by stores (@ store execute) to detect load ordering violation
- ❑ **Both together**
  - Allow aggressive load scheduling
    - Stores don't constrain load execution
  - Help us improve performance significantly

(163)

163

## Multi-threading (MT)


- ❑ Even moderate static superscalars (e.g., 4-way) are not fully utilized
  - Average sustained IPC: 1.5–2 → < 50% utilization due to
    - Mispredicted branches (branch frequency?)
    - Cache misses, especially L1 (very frequent)
    - Data dependences, load-use data hazards
- ❑ Multi-threading (MT) to the rescue
  - Improve utilization of datapath components by multiplexing multiple (process) threads on single datapath
  - If one thread cannot fully utilize the datapath, maybe 2 or 4 (or 100) can

(165)


165

## Multithreading Example

- ❑ Time evolution of issue slot



- ❑ # cycles? # wasted cycle slots?



Multithreaded Static SS

- ❑ Fill in with instructions from other threads – in this example we have 2 threads and change threads every cycle

- Completely removes load-use hazard empty slots
- Takes longer for the “red” thread to finish
  - With more threads, would take even longer
- Still have some noop slots (so wasted performance – stay tuned)

(166)

166



## Latency vs throughput

### □ MT trades (single-thread) latency for throughput

- Sharing processor degrades latency of individual threads
- + But improves aggregate latency of both threads
- + Improves utilization

### □ Example

- Thread A: individual latency=10s, latency with thread B=15s
- Thread B: individual latency=20s, latency with thread A=25s
- Sequential latency (first A then B or vice versa): 30s
- Parallel latency (A and B simultaneously): 25s
- MT slows each thread by 5s
- + But improves total latency by 5s

### □ Different workloads have different types of parallelism

- SpecFP has lots of ILP; i.e., can make use of an 8-wide machine
- Server workloads have TLP (Thread Level Parallelism), i.e., have multiple threads that can run in parallel

167

## Alternative Multithreaded Implementations

### □ MT trades (single-thread) latency for throughput

- Sharing the datapath degrades the latency of individual threads, but improves the aggregate latency of both threads
- And, it improves utilization of the datapath hardware

### □ Main questions: thread scheduling policy and pipeline partitioning

- When to switch from one thread to another?
- How exactly do threads share the pipelined datapath itself?

### □ Choices depends on what kind of latencies you want to tolerate and how much single thread performance you are willing to sacrifice

- Coarse-grain multithreading (CGMT) ←
- Fine-grain multithreading (FGMT)
- Simultaneous multithreading (SMT)


(168)

168

## The standard multithreading pictures

□ Time evolution of issue slots

● Color = thread



(169)

169

## Coarse-Grain MultiThreading (CGMT)

+ Sacrifices very little single thread performance (of one thread)

- Tolerates only long latencies (e.g., L2 misses)


□ Thread scheduling policy

- Designate a "preferred" thread (e.g., thread A)
- Switch to thread B on thread A L2 miss
- Switch back to A when A L2 miss returns

□ Pipeline partitioning

- None, flush on switch
- So can't tolerate very short latencies
  - Need short in-order pipeline for good performance

□ Example: IBM Northstar/Pulsar

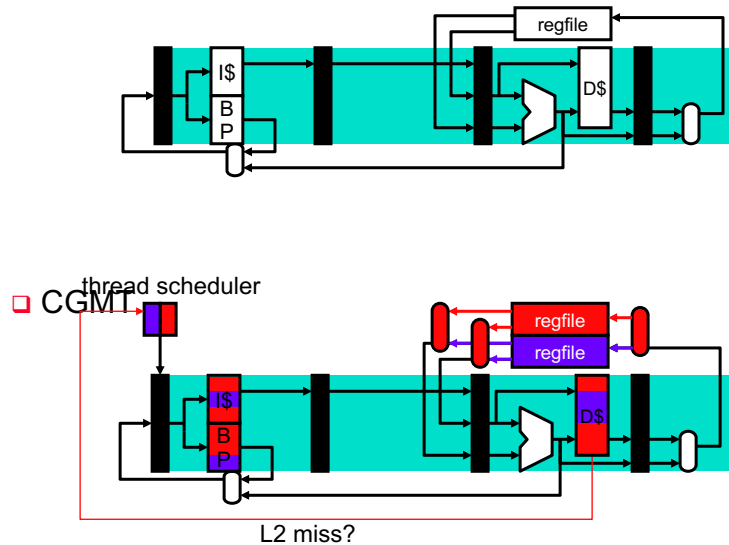


CGMT

(170)

170

## Coarse-Grain MultiThreaded architecture



(171)

171

## Alternative Multithreaded Implementations

- ❑ MT trades (single-thread) latency for throughput
  - Sharing the datapath degrades the **latency** of individual threads, but improves the aggregate latency of both threads
  - And, it improves **utilization** of the datapath hardware
- ❑ Main questions: **thread scheduling policy and pipeline partitioning**
  - When to switch from one thread to another?
  - How exactly do threads share the pipelined datapath itself?
- ❑ Choices depends on what kind of latencies you want to tolerate and how much single thread performance you are willing to sacrifice
  - Coarse-grain multithreading (**CGMT**)
  - Fine-grain multithreading (**FGMT**) ←
  - Simultaneous multithreading (**SMT**)

(172)

172

## Fine-Grain MultiThreading (FGMT)

- Sacrifices significant single thread performance
- + Tolerates latencies (e.g., load-use hazards, L1 misses, mispredicted branches, etc.)
- ❑ Thread scheduling policy
  - Switch threads every cycle (round-robin, can skip stalled threads)
- ❑ Pipeline partitioning
  - Dynamic, no pipeline flushing between threads
- Need a lot of threads
- ❑ Extreme example: Denelcor HEP
  - So many threads (100+), it didn't even need caches
  - Targeted for DoD, not successful commercially
  - [http://en.wikipedia.org/wiki/Heterogeneous\\_Element\\_Processor](http://en.wikipedia.org/wiki/Heterogeneous_Element_Processor)
- ❑ Sun's UltraSPARC T1 (Niagara)
  - Many threads → many RF [http://en.wikipedia.org/wiki/UltraSPARC\\_T1](http://en.wikipedia.org/wiki/UltraSPARC_T1)



FGMT

(173)

173

## FGMT Sharing Implementations Issues

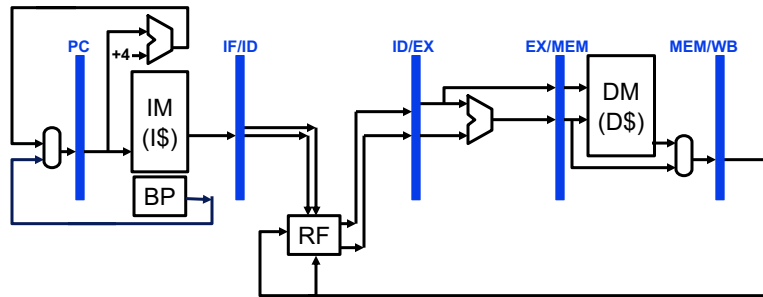
- ❑ How do multiple threads share a single datapath?
  - Different sharing mechanisms for different kinds of structures, depending on what kind of **state** the structure stores
- ❑ **No state**: ALUs
  - So, can be dynamically shared
- ❑ **Persistent hard state (aka thread "context")**: PC, RFile
  - So must be **replicated**
- ❑ **Persistent soft state**: caches, TLBs, branch prediction structures (BTB, BHT)
  - Dynamically partitioned (like on a multi-programmed uni-processor)
    - TLBs need thread ids, caches/branch prediction table (BHT) don't
- ❑ **Transient state**: pipeline latches
  - Must be partitioned ... somehow

(174)

174

## FGMT Datapath (Single Issue)

- What do we have to add to our datapath to support FGMT?

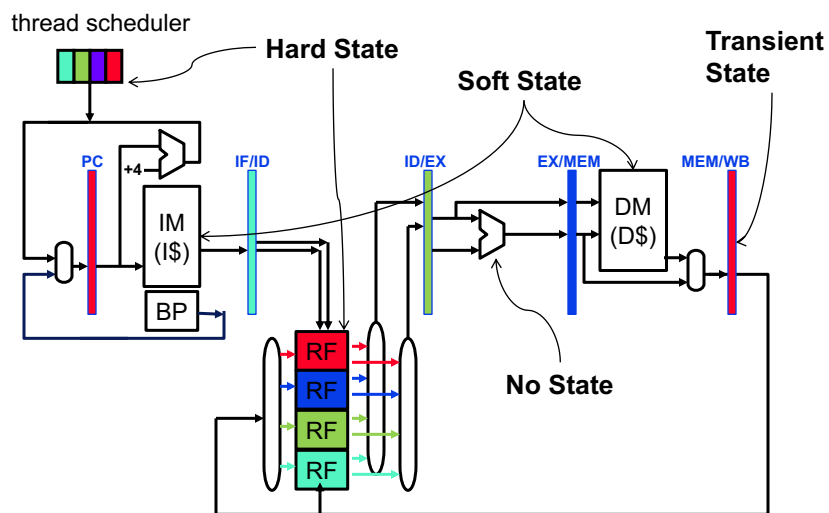


(175)

175

## FGMT Datapath (Single Issue)

- What do we have to add to our datapath to support FGMT?

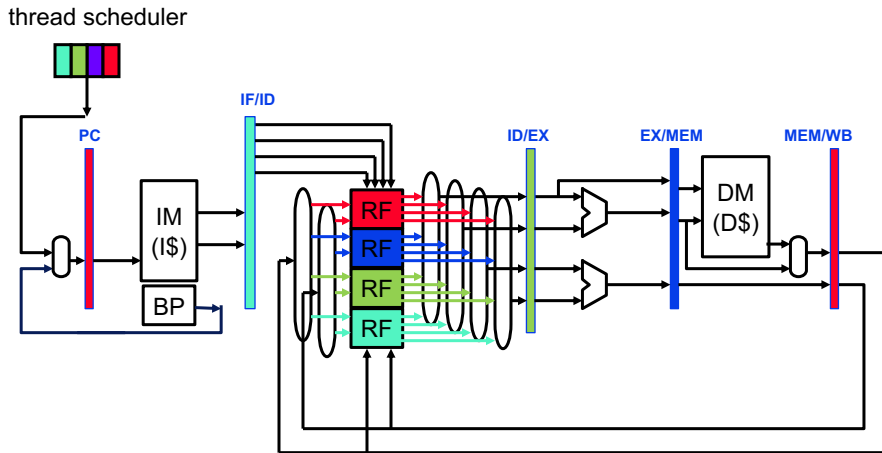


(176)

176

## FGMT Datapath (2-Way Issue)

- ❑ What do we have to add to our datapath to support FGMT?

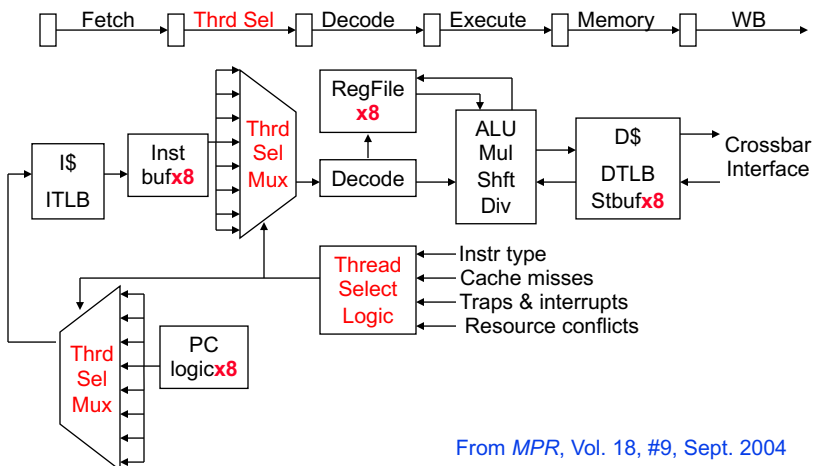


(177)

177

## Sun Niagara's FGMT Integer pipeline

- ❑ Cores are simple (single-issue, 6 stage, no branch prediction), small, and power-efficient



From *MPR*, Vol. 18, #9, Sept. 2004

(178)

178

- 8 SPARC FGMT datapath cores



179

- Time evolution of issue slots
  - Color = thread





## Simultaneous Multithreading (§3.10)

- Key idea

Issue multiple instructions from multiple threads each cycle

- Features

- Fully exploit thread-level parallelism and instruction-level parallelism.
- Better Performance for
  - Mix of independent programs
  - Programs that are parallelizable

(181)

181

## Simultaneous MultiThreading (SMT)

- What can issue instr's from multiple threads in one cycle?

- Same thing that issues instr's from multiple parts of same thread ...
- ...out-of-order execution !!

- Simultaneous multithreading (SMT): **OOO + FGMT**

- Most common implementation of multithreading!
- Aka (by Intel) "hyper-threading"
- Once instr's are renamed, issuer doesn't care which thread they come from (well, for non-loads at least)
- Some examples
  - IBM Power5: 4-way, 2 threads; IBM Power7: 4-way, 4 threads
  - Intel Pentium4: 3-way, 2 threads; Intel Core i7: 4-way, 2 threads
  - AMD Bulldozer: 4-way, 2 threads
  - Alpha 21464: 8-way issue, 4 threads (canceled)
  - Notice a pattern?  $\#threads (T) * 2 = \#issue\ width (N)$

(182)

182






## Multithreading Speed-Ups on the i7 (SMT)

Multithreading allows multiple threads to share the functional units of a single processor in an overlapping fashion

Duplicates only private state – not everything



i)

185

## Multithreading vs Multicore

- ❑ If you wanted to run multiple threads would you build a
  - A multicore: multiple separate pipelines?
  - A multithreaded processor: a single larger pipeline?
- ❑ **Both will get you throughput on multiple threads**
  - A multicore core will be simpler, possibly faster clock
    - Multicore is mainly a TLP (thread-level parallelism) engine
  - SMT will get you better performance (IPC) on a single thread
    - SMT is basically an ILP engine that converts TLP to ILP
- ❑ **Do both**
  - Intel's Sandy (Ivy) Bridge and Haswell, IBM's Power7 & 8
  - 4 to 8 OOO 4-way cores each of which supports 2 to 4 threads (SMT)
  - Private L1 and L2 caches, shared L3 cache
  - 3+ GHz clock rate
  - Graphics processors

(186)

186

### Check Yourself

- ☐ Branch prediction
- ☐ Multiple issue
- ☐ VLIW
- ☐ Superscalar
- ☐ Static scheduling
- ☐ Dynamic scheduling
- ☐ In-order execution
- ☐ Out-of-order (OOO) execution
- ☐ Speculation
- ☐ Reorder buffer
- ☐ Register renaming

(187)